

JANUARY 2007						
M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

FEBRUARY 2007						
M	T	W	T	F	S	S
5	6	7	8	9	10	11
12	13	14	15	16	17	18
19	20	21	22	23	24	25
26	27	28				

Schlageter proposed an environment to King Robinson method whereby transaction always proceeds without a validation check and thus without the risk of restarts. DECEMBER 2017 SATURDAY WK - 50 - 343-022

09

8.00 Validation Protocol: → In all concurrency control techniques (locking & timestamping), a certain degree of checking is done before database operation can be executed. For example in a locking, a check is done to determine whether item being accessed is locked. In timestamp ordering, the transaction timestamp is checked against read and write timestamp of the item. Such checking represents overhead during transaction execution, with effect of slowing down the transaction and thereby reducing the concurrency. According to King and Robinson; no checking is done while transaction is executed. It is optimistic because the likelihood of two clients transactions accessing the same object is low. Transactions are allowed to proceed as though there is no possibility of conflict with other transaction until client completes its task. When conflict arises, some transaction is generally aborted and will need to be restarted by the client.

5.00  
6.00  
Evening  
As Validation & update phase are generally short in duration compared to working phase, only one transaction may be in validation and update phase at one time.

In validation phase for transaction  $T_i$ , the protocol checks that  $T_i$  doesn't interfere with any committed transaction or with any other transaction currently in their validation phase.

The validation phase for  $T_i$  checks that for each such transaction  $T_j$  that is either committed or is in its validation phase, one of the following conditions holds:

- 1) Transaction  $T_j$  completes its write phase before  $T_i$  starts its read phase.  $\text{Finish}(T_j) < \text{Start}(T_i)$
- 2) Set of data items written by  $T_i$  doesn't intersect with set of data items read by  $T_j$  and  $T_i$  completes its write phase before  $T_j$  starts its validation phase.  $\text{Writes}(T_i) \cap \text{Reads}(T_j) = \emptyset$



verening

There are 3 phases in validation protocol:  $\rightarrow$

1) Read phase:  $\rightarrow$  Transaction read values from committed data items from database. Updates are applied only to local copies of data items.

2) Validation phase:  $\rightarrow$  Checking is performed to ensure serializability will not be violated if the transaction update is applied to database.

Notes

3) Write phase:  $\rightarrow$  If validation phase is successful, update are applied to database otherwise updates are discarded and transaction restarted.

The idea behind this technique is to do all checks at once; hence transaction execution proceeds with minimum overhead until validation phase is reached.

don't overlap.

Multiversion timestamp:  $\Rightarrow$  It was introduced by Reed in 1983. A list of old committed versions as well as tentative versions is kept for each object. This list represents the history of the values of the object. The benefit of using multiple versions is that read operations that arrived too late need not be rejected.

When a transaction requires access to an item an appropriate version is chosen to maintain serializability of the currently executing schedule if possible.

An obvious drawback of this technique is that more storage is needed to maintain multiple versions of the database items. However, older versions may have to be maintained anyway - for example for recovery purposes. They are mainly used in temporal database which keeps track of all changes and the time at which they occurred.

Each version has following fields:  $\Rightarrow$

- i) Content of  $X$ .
- ii) read timestamp  $\Rightarrow$  largest of all timestamps of transaction that have successfully read version  $X$ .
- iii) write timestamp  $\Rightarrow$  Timestamp of transaction that wrote value of  $X$ .

(with timestamp  $T$ )

Whenever a transaction  $T_i$  is allowed to execute a write operation, a new version  $X_{k+1}$  of item  $X$  is created, with readtimestamp and writetimestamp set to  $TS(T_i)$ . Correspondingly, when a transaction  $T_i$  is allowed to read the value of version  $X_i$ , the readtimestamp is set to

JANUARY 2007						
M	T	W	T	F	S	S
1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

FEBRUARY 2007						
M	T	W	T	F	S	S
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

When a version  $X_t$  has write time  $t$  such that a later transaction  $T$  has a timestamp less than  $t$ , then we may delete version of  $X$  previous to  $X_t$ . DECEMBER 2007, TUESDAY

WK-51-946-019

12

0.00 Larger of current read timestamp and  $TS(T)$ .

9.00 To ensure serializability, following rules around.  
a) If Transaction  $T$  issues read( $x_i$ ), find the version of  $X$  that has highest write timestamp of all versions of  $X$  that is equal or less than  $TS(T)$  and  $readTS(x_i) > TS(T)$  then return the value of  $x_i$  to transaction  $T$  and set the value of  $readTS(x_i)$  to larger of  $TS(T)$  and current  $readTS(x_i)$ .

b) If transaction  $T$  issues write( $x_i$ ) and  $TS(T) < ReadTimestamp(x_i)$ , then about  $T$  other create a new version of  $X$  with  $r(T)=w(T)=TS(T)$ .

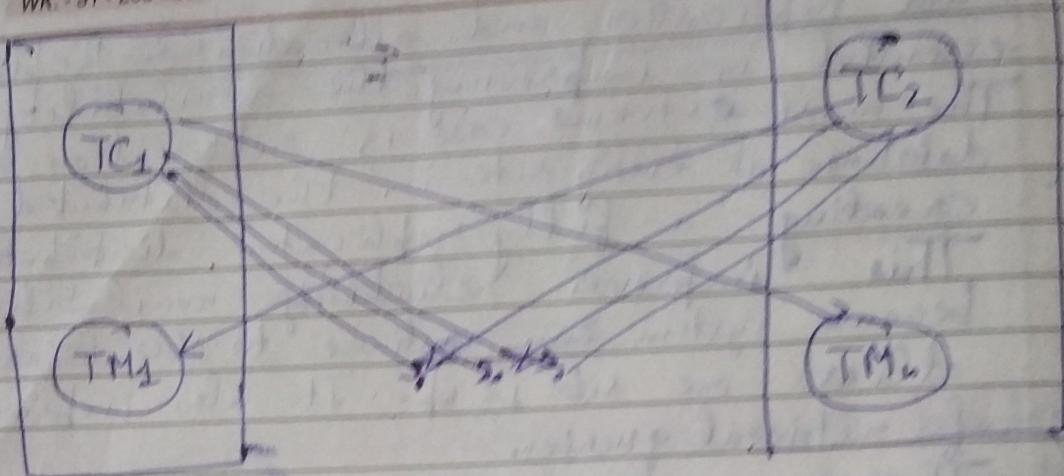
~~It has~~

4.00 ~~\* Multiversion 2 Phase locking.~~ This protocol differentiates read only transactions and update transactions. Update transactions perform rigorous 2 PL.  
5.00 In multiversion 2 PL, we perform allow other transactions  $T'$  to read an item  $X$  while a single transaction  $T$  holds a write lock on  $X$ . This is accomplished by allowing two versions for each item  $X$ ; one version must have always been by some committed transaction. The second version  $X'$  is created when a transaction acquires write lock on item. Other transactions can continue to read ~~and~~ the committed version of  $X$  while  $T$  holds write lock.

Notes  
Once  $T$  is ready to commit, it must obtain certified lock on all items that it currently holds write locks on before it can commit. All read locks on these items be released <sup>in order to</sup> ~~by~~ use certified lock. Once certified lock <sup>and</sup> acquired ~~and~~  $x$  is set to value of version  $X'$  and  $X$  is deleted.

0000/11/26 To ensure atomicity, all sites on which transaction is executed must agree on the final outcome. If any participant executes a transaction, it must either commit it at all sites or must abort it at all sites.

MONDAY				TUESDAY				WEDNESDAY				THURSDAY				FRIDAY														
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	
31	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30
31	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30



TC is responsible for

- Starting execution of the transaction.
- Breaking transaction into a number of subtransactions and distributing these subtransactions to appropriate sites for execution.
- Coordinating the outcome of the transaction, which may result in the transaction being committed at all sites or aborted at all sites.

Basic failure of dbms:

- Failure of site
- Loss of message
- Failure of communication link
- N/W Partition

Solutions:

Redo log

Retransmission

Recovery

Check for consistency

Two phase Commit protocol:

Goal of 1st phase  $\rightarrow$  To reach a common decision  
 1. 2nd,  $\rightarrow$  To implement the decision

1st phase  $\rightarrow$  The coordinator is responsible for taking the final commit or abort decision. Each participant

AUGUST 2006						
M	T	W	T	F	S	S
1	2	3	4	5	6	
7	8	9	10	11	12	13
14	15	16	17	18	19	20
21	22	23	24	25	26	27
28	29	30	31			

SEPTEMBER 2006						
M	T	W	T	F	S	S
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29	30	

JULY 2006

TUESDAY  
WK. - 31 • 206-159

25

8.00 corresponds to a subtransaction which has performed some write action i.e. it is responsible for performing write action at its local database.

10.00 Phase I: → Here, the coordinator asks all the participants to prepare for commitment; each participant answers Ready if it is ready to commit. Before sending the message, the coordinator records on stable storage a log record of <prepareT>. On receiving such a message, the transaction mgr at that site determines whether it is willing to commit its portion of T. If the answer is no, it adds record <no T> to the log and then responds by sending abort T message to Ci. If the answer is yes, it adds a record <readyT> to the log and then forces the log onto stable storage. Then the TM then replies with a ready T message to Ci.

5.00 The coordinator decides whether to commit or abort as a result of the answer it received from participants. If all the participants answer ready, it decides to commit. If some participant has answered Abort or has not yet answered when the timeout expires Evening it decides to abort the transaction. Depending on verdict, the fate of the transaction has been sealed. either a record <commitT> or a record <abortT> is added to the log. Afterwards, the coordinator sends either a commit T or abort T message to all participants. When site receives message it records the message in the log.

Finally, all participants send a final acknowledgement message to the coordinator and perform the actions required for committing or aborting subtransaction when the Coordinator has received an ACK message from all participants, it writes a log record called the COMMIT from all participants, it writes a log record called

JULY 2006						
M	T	W	T	F	S	S
31			1	2		
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30

AUGUST 2006						
M	T	W	T	F	S	S
	1	2	3	4	5	6
7	8	9	10	11	12	13
14	15	16	17	18	19	20
21	22	23	24	25	26	27
28	29	30	31			

JUNE 2006

08

THURSDAY

WK. - 24 • 159-206

8.00

Multiple Granularity → A database item could

- a record ✓
- a field value ✓
- disk block ✓
- whole file ✓
- whole database ✓

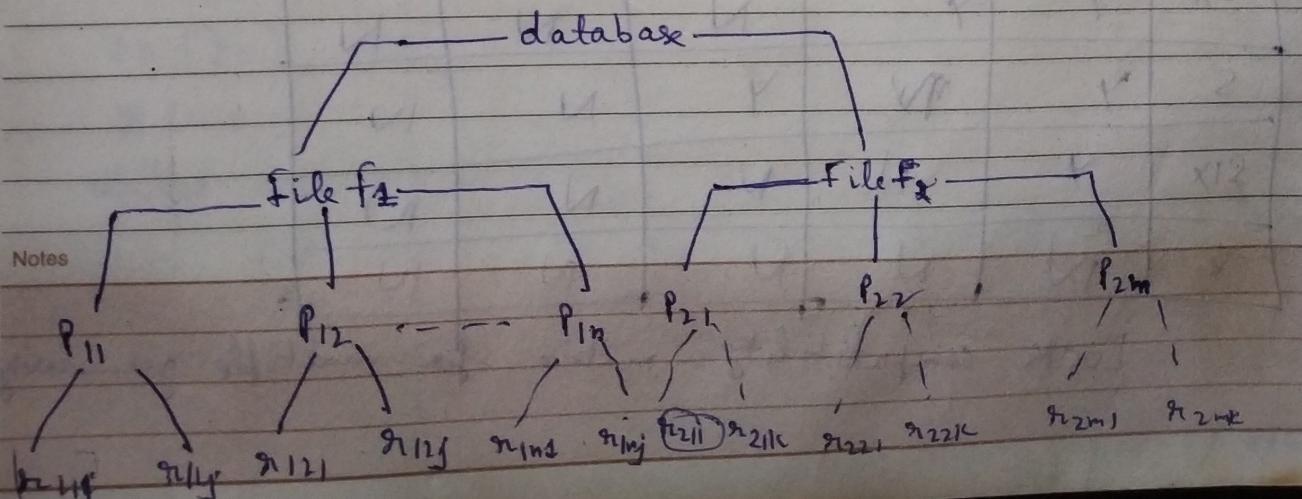
The granularity can affect the performance of concurrency control and recovery. The size of data items is known as data item granularity. Fine granularity refers to small item sizes whereas coarse granularity refers to large item sizes. Larger the data item size, lower is the degree of concurrency permitted.

Smaller the size, more is the no. of active holds causing higher overhead. Also, more storage space required for lock table.

So, what is the best item size? Well, it depends on the types of transactions involved.

Case:

- If transaction accesses a small no. of records, then we can make single record as data item.
- If transaction involves many records in the same file, it may better to have block or file granularity.



2006 JUNE

09

FRIDAY

WK. 24 \* 160-208

MAY 2006							JUNE 2006						
M	T	W	T	F	S	S	M	T	W	T	F	S	S
1	2	3	4	5	6	7	8	9	10	11	12	13	14
15	16	17	18	19	20	21	22	23	24	25	26	27	28
29	30	31					26	27	28	29	30		

8.00 The multiple granularity protocol requires that  
9.00 lock be acquired in top-down (root to leaf) order,  
whereas locks must be released in bottom-up (leaf-to-root) order.

10.00 Rules:-

- 1) Adhere to lock compatibility matrix.
- 2) Root of tree must be locked first and can be locked in any mode.
- 3) A node  $\varnothing$  can be locked by  $T_i$  in S or IS mode only if the parent of  $\varnothing$  is currently locked by  $T_i$  in either IX or IS mode.
- 4) A node  $\varnothing$  can be locked by  $T_i$  in X, IX or SIX only if the parent of  $\varnothing$  is currently locked by  $T_i$  in either IX or SIX mode.
- 5)  $T_i$  can lock a node only if it has not previously unlocked any node i.e.  $T_i$  is two-phase.
- 6)  $T_i$  can unlock a node  $\varnothing$  if none of children of  $\varnothing$  are currently locked by  $T_i$ .

	IS	IX	S	SIX	X
IS	Y	Y	Y	Y	No
IX	Y	Y	N	N	N
S	Y	Y	Y	N	N
SIX	Y	N	N	N	N
X	N	N	N	N	N

lock compatibility matrix for multiple granularity



REDMI NOTE 5 PRO  
MI DUAL CAMERA

JULY 2020

1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

AUGUST 2020

1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30	31				

JUNE 2020

SATURDAY  
08.06.2020

10

\* The idea behind intention locks is for a transaction to indicate, along the path from root to the desired node what type of lock (shared or exclusive) it will require of the nodes descendants.

- IS → shared lock will be requested on some descendant node.
- IX → exclusive lock will be requested on some descendant node.
- SIS → current node is locked in shared mode but exclusive lock will be requested on some descendant node.

a) Update record  $r_{11} \& r_{21}$       b) update records on page  $p_{12}$       c) records in another file

T1

IX (db)

IX (f1)

IX (p11)

X (r11)

IX (f2)

IX (p21)

X (r21)

unlock  $r_{11}, r_{21}$ ,  $p_{21}$ 

f2

r11

T2

IX (db)

IX (f1), IX (p12)

IS (db)

IS (f1)

IS (p11)

S (r11)

SUNDAY 11

S (f1)

unlock ( $p_{12}, f_1, db$ )unlock  $r_{11}, p_{11}, db, f1$